
EXERCISE ON NÜSKEN-ZIEGLER ALGORITHM

1. MULTIPOINT EVALUATION OF BIVARIATE POLYNOMIALS

We consider the following two algorithmic problems:

BivMPEval: Given parameters $n, d_x, d_y \in \mathbb{Z}_{>0}$, given n points $(a_1, b_1), \dots, (a_n, b_n)$ in \mathbb{K}^2 with the a_i 's pairwise distinct, given a polynomial $F(x, y) \in \mathbb{K}[x, y]$ with $\deg_x(F) < d_x$ and $\deg_y(F) < d_y$, compute the evaluations $F(a_1, b_1), \dots, F(a_n, b_n)$.

This is bivariate multipoint evaluation, with the restriction that the x -coordinates are distinct. At the price of some “mild randomization”, one can reduce to this situation from the general case of n distinct points; this reduction will not be discussed here.

BivModComp: Given parameters $n, d_x, d_y \in \mathbb{Z}_{>0}$, given a polynomial $F(x, y) \in \mathbb{K}[x, y]$ with $\deg_x(F) < d_x$ and $\deg_y(F) < d_y$, given polynomials $P, Q \in \mathbb{K}[x]$ with $\deg(P) < \deg(Q) = n$, compute the modular composition $F(x, P(x)) \bmod Q(x)$.

This is one type of bivariate extension of the modular composition of univariate polynomials.

- (1) Give a complexity bound for a naive algorithm for **BivMPEval**, which performs n evaluations of F at a single point.
- (2) Exploiting the polynomials $Q = \prod_{1 \leq i \leq n} (x - a_i)$ and $P \in \mathbb{K}[x]_{<n}$ such that $P(a_i) = b_i$ for $1 \leq i \leq n$, prove that: from an algorithm \mathcal{A} for **BivModComp** with complexity $\mathcal{C}(n, d_x, d_y)$, one can derive an algorithm \mathcal{B} for **BivMPEval** whose complexity is bounded by $\mathcal{C}(n, d_x, d_y) + O(M(n) \log(n))$.
- (3) Give an algorithm for **BivModComp** based on the writing $F = \sum_{j < d_y} F_j(x) y^j$ and on “naive” evaluation at $y = P \bmod Q$. Show that its complexity is in $O(d_y M(d_x + d_y M(n)))$, and give a range for (d_x, n) for which this is quasi-linear in the input size.

From here on, we focus on the case $d_x \in O(n)$. Our goal is to obtain a better complexity bound for **BivModComp** (thus also for **BivMPEval**) than the one $O(d_y M(n))$ from Question (3).

- (4) If $d_y \in O(n)$ and $d_x = 1$, do you know a better complexity bound for **BivModComp**?
- (5) For simplicity, we now also assume (in addition to $d_x \in O(n)$) that d_y is a square $d_y = \delta^2$ with $\delta \in \mathbb{Z}_{>0}$. Writing $F = \sum_{j < \delta} (\sum_{i < \delta} F_{i+\delta j}(x) y^i) y^{\delta j}$, give an algorithm for **BivModComp** which exploits polynomial matrix multiplication.
[Hint: compute $\sum_{i < \delta} F_{i+\delta j}(x) P^i \bmod Q$ simultaneously for all j via the multiplication of a $\delta \times \delta$ univariate matrix of degree $< d_x$ by a $\delta \times 1$ univariate vector of degree $< n$.]
- (6) Recall how to perform the above matrix-vector product using $O(\delta^\omega M(\frac{n}{\delta} + d_x))$ operations in \mathbb{K} . Deduce that the algorithm of Question (5) costs $O(\delta M(n) + \delta^\omega M(\frac{n}{\delta} + d_x))$.
- (7) As soon as $d_x \in O(\frac{n}{\delta})$ (in particular, in the frequent case $d_x d_y \in \Theta(n)$), the above bound is within $O(\delta^{\omega-1} M(n))$. How does this compare to $O(d_y M(n))$?

Solution:

- (1) The evaluation of F at a single point (a_i, b_i) costs $O(d_x d_y)$ operations in \mathbb{K} . Evaluating naively at each point independently will thus cost $O(n d_x d_y)$ operations in \mathbb{K} .
- (2) Define the *univariate* polynomial $R = F(x, P) \bmod Q$. Then $R = F(x, P(x)) + A(x)Q(x)$ for some $A \in \mathbb{K}[x]$. Since Q has roots a_1, \dots, a_n , we get, for $1 \leq i \leq n$, $R(a_i) = F(a_i, P(a_i)) + A(a_i)Q(a_i) = F(a_i, b_i)$. Therefore an algorithm \mathcal{B} to solve **BivMPEval** is
 - compute P and Q from the points (subproduct tree & Lagrange interpolation);
 - use algorithm \mathcal{A} to solve **BivModComp** on input P, Q, F , which provides the polynomial R above;
 - compute $R(a_1), \dots, R(a_n)$ (univariate multipoint evaluation).

The complexity of the first and third steps is $O(M(n) \log(n))$, hence the sought overall bound $\mathcal{C}(n, d_x, d_y) + O(M(n) \log(n))$.

(3) Writing $F = \sum_{j < d_y} F_j(x)y^j$, we obtain

$$F(x, P(x)) \bmod Q(x) = \left(\sum_{j < d_y} F_j(x) (P(x)^j \bmod Q(x)) \right) \bmod Q(x).$$

This formula leads straightforwardly to solving **BivModComp** in $O(d_y \mathbf{M}(d_x) + d_y \mathbf{M}(n))$ operations in \mathbb{K} . The above complexity is quasi-linear in the size of the input when $n \in O(d_x)$.

- (4) In this case, Paterson and Stockmeyer's (or Brent and Kung's) baby step giant step algorithm uses $O(n^{\frac{\omega+1}{2}} + n^{\frac{1}{2}} \mathbf{M}(n))$ operations in \mathbb{K} .
- (5) Rewriting indices, we have

$$F = \sum_{i,j < \delta} F_{i+\delta j}(x)y^{i+\delta j} = \sum_{j < \delta} \left(\sum_{i < \delta} F_{i+\delta j}(x)y^i \right) y^{\delta j} = \sum_{j < \delta} \hat{F}_j(x, y)y^{\delta j}$$

where $\hat{F}_j(x, y) = \sum_{i < \delta} F_{i+\delta j}(x)y^i$, for $j < \delta$. We can solve **BivModComp** by solving several instances of it (each with smaller y -degree) with the polynomials $\hat{F}_j(x, y)$:

- for $j < \delta$, compute $R_j = \hat{F}_j(x, P(x)) \bmod Q(x)$ [δ instances of **BivModComp**]
- for $j < \delta$, compute $P_j = P^{\delta j} \bmod Q$ [total cost: $O(\delta \mathbf{M}(n))$]
- compute and return $\sum_{j < \delta} (R_j P_j \bmod Q)$ [cost $O(\delta \mathbf{M}(n))$]

The complexity gain is obtained by realizing the computations of the R_j 's simultaneously, using polynomial matrix multiplication:

$$\begin{bmatrix} R_0 \\ R_1 \\ \vdots \\ R_{\delta-1} \end{bmatrix} = \begin{bmatrix} F_0 & F_1 & \cdots & F_{\delta-1} \\ F_\delta & F_{1+\delta} & \cdots & F_{2\delta-1} \\ \vdots & \vdots & \cdots & \vdots \\ F_{(\delta-1)\delta} & F_{(\delta-1)\delta+1} & \cdots & F_{\delta^2-1} \end{bmatrix} \begin{bmatrix} 1 \\ P \\ P^2 \bmod Q \\ \vdots \\ P^{\delta-1} \bmod Q \end{bmatrix} \bmod Q$$

This can be split into several steps:

- compute the powers $P^i \bmod Q$ for $i < \delta$ [total $O(\delta \mathbf{M}(n))$]
- compute the matrix-vector product, efficiently by expanding the vector into a $\delta \times \delta$ matrix of degree less than n/δ [total $O(\delta^\omega \mathbf{M}(\frac{n}{\delta} + d_x))$]
- the previous step has yielded δ polynomials of degree less than $n + d_x \in O(n)$ [total $O(\delta \mathbf{M}(n))$].
 \rightarrow reduce them mod Q

(6) Complexity bounds are given in the answer just above.

(7) Yes, since $O(\delta^{\omega-1} \mathbf{M}(n)) = O(d_y^{\frac{\omega-1}{2}} \mathbf{M}(n))$, and $\frac{\omega-1}{2} < 0.7$ with the best known value of $\omega < 2.4$.